

3.2 Example : Linear regression

In this example, we will use the dataset “Advertising and Sales”, which is available at www.kaggle.com. The datasets includes 4 features and 200 instances. The 3 features are “TV”, “radio”, “newspaper”, which refer to the amount of resources that are allocated in each media. The 4th feature is the target feature “sales”, which refers to the total sales number. For the purpose of this chapter, we will assume that the only feature that has strong impact on the target feature is “TV”. Thus, we ignore the other two features. In addition, we will use 160 out of the 200 instances. For this example we will use the programming language “python”.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
```

```
In [2]: df=pd.read_csv('Advertising.csv')
df.drop('Unnamed: 0',inplace=True, axis=1)
df
```

Out[2]:

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	9.7
197	177.0	9.3	6.4	12.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	13.4


200 rows × 4 columns

First, we will handle our data because they should have a specific form, before they are used as an input to our model. Then we will fit the model to our data and we will plot the results.

```
In [3]: x = df[:160]['TV'].values
y = df[:160]['sales'].values
x=x.reshape(160,1)
y=y.reshape(160,1)
```

```
In [4]: lr = LinearRegression().fit(x,y)
y_pred=lr.predict(x)
plt.scatter(x,y)
plt.xlabel('TV')
plt.ylabel('sales')
plt.plot(x,y_pred,color='red')
```

```
Out[4]: [<matplotlib.lines.Line2D at 0x2289b1c5df0>]
```



```
In [5]: x1 = sm.add_constant(x)
est = sm.OLS(y, x1)
est2 = est.fit()
print(est2.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.643
Model:                  OLS    Adj. R-squared:            0.641
Method:                 Least Squares    F-statistic:        284.6
Date:                   Mon, 29 Jan 2024    Prob (F-statistic):  3.58e-37
Time:                   10:31:23    Log-Likelihood:     -408.80
No. Observations:       160    AIC:                821.6
Df Residuals:           158    BIC:                827.7
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	7.0688	0.483	14.634	0.000	6.115	8.023
x1	0.0489	0.003	16.871	0.000	0.043	0.055

```

=====
Omnibus:                 1.575    Durbin-Watson:       1.931
Prob(Omnibus):            0.455    Jarque-Bera (JB):     1.414
Skew:                    -0.230    Prob(JB):             0.493
Kurtosis:                 3.002    Cond. No.             325.
=====

```

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Evaluation

```
In [6]: x_test = df[160:]['TV'].values
y_test = df[160:]['sales'].values
x_test = x_test.reshape(40,1)
y_test = y_test.reshape(40,1)
```

```
In [7]: y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test)
plt.xlabel('TV')
plt.ylabel('sales')
plt.plot(x_test,y_pred,color='red')
```

Out[7]: [



```
In [8]: from sklearn import metrics
MSE = metrics.mean_squared_error(y_test,y_pred)
print('MSE =', round(MSE,2))
MAE = metrics.mean_absolute_error(y_test,y_pred)
print('MAE =', round(MAE,2))
RMSE = MSE**(1/2)
print('RMSE =', round(RMSE,2))
R2 = metrics.r2_score(y_test,y_pred)
print('R-squared =', round(R2,2))
```

```
MSE = 14.13
MAE = 3.08
RMSE = 3.76
R-squared = 0.47
```